

```
In [9]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import fetch_openml
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.decomposition import PCA
from sklearn.manifold import TSNE
```

```
In [16]: # Dataset Loading
mnist = fetch_openml('mnist_784', version=1, as_frame=False)
X, y = mnist.data, mnist.target.astype(np.int8)
```

```
In [22]: X = X / 255.0 # Normalize pixel values
# Reduce dimensionality using PCA
pca = PCA(n_components=50) # You can try 30-100 and tune this
X_pca = pca.fit_transform(X)
```

```
In [31]: # Step 3: Model Development (Using GaussianNB as approximation to Bayes' Decision Rule)
model = GaussianNB()

# Step 4: Training and Testing
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
```

```
In [32]: # Step 5: Evaluation
print("\nEvaluation Metrics:")
print(f"Accuracy: {accuracy_score(y_test, y_pred) * 100:.2f}%")
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

Evaluation Metrics:


Accuracy: 55.16%

Classification Report:

	precision	recall	f1-score	support
0	0.69	0.91	0.78	1343
1	0.81	0.95	0.88	1600
2	0.85	0.30	0.44	1380
3	0.74	0.32	0.45	1433
4	0.83	0.13	0.23	1295
5	0.61	0.04	0.08	1273
6	0.64	0.94	0.76	1396
7	0.91	0.28	0.42	1503
8	0.29	0.60	0.39	1357
9	0.37	0.94	0.53	1420
accuracy			0.55	14000
macro avg	0.67	0.54	0.50	14000
weighted avg	0.68	0.55	0.51	14000

```
In [25]: conf_mat = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(10, 7))
sns.heatmap(conf_mat, annot=True, fmt='d', cmap='Blues')
```

```
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("True")
plt.show()
```

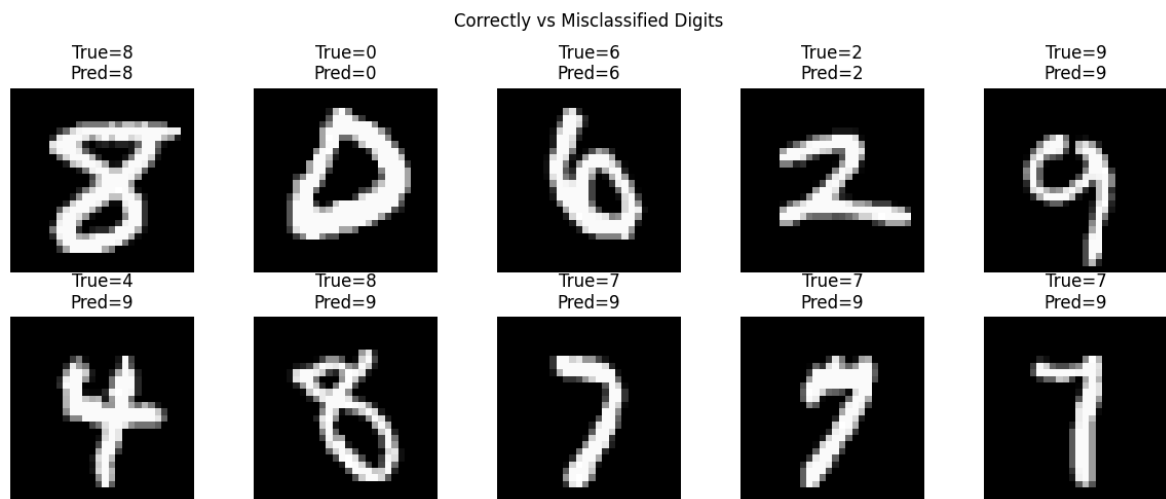


```
In [29]: # Visualization - Correct vs Misclassified
correct = np.where(y_pred == y_test)[0]
incorrect = np.where(y_pred != y_test)[0]

plt.figure(figsize=(12, 5))
for i, idx in enumerate(correct[:5]):
    plt.subplot(2, 5, i + 1)
    plt.imshow(X_test[idx].reshape(28, 28), cmap='gray')
    plt.title(f"True={y_test[idx]}\nPred={y_pred[idx]}")
    plt.axis('off')

for i, idx in enumerate(incorrect[:5]):
    plt.subplot(2, 5, i + 6)
    plt.imshow(X_test[idx].reshape(28, 28), cmap='gray')
    plt.title(f"True={y_test[idx]}\nPred={y_pred[idx]}")
    plt.axis('off')


plt.suptitle("Correctly vs Misclassified Digits")
plt.tight_layout()
plt.show()
```



```
In [27]: # Visualize Decision Boundaries with PCA (2D)
print("Reducing dimensionality for visualization...")
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_test)

plt.figure(figsize=(10, 7))
scatter = plt.scatter(X_pca[:, 0], X_pca[:, 1], c=y_pred, cmap='tab10', alpha=0.1)
plt.legend(*scatter.legend_elements(), title="Predicted")
plt.title("2D PCA Projection Colored by Predicted Labels")
plt.xlabel("PCA 1")
plt.ylabel("PCA 2")
plt.grid(True)
plt.show()
```

Reducing dimensionality for visualization...



In []: